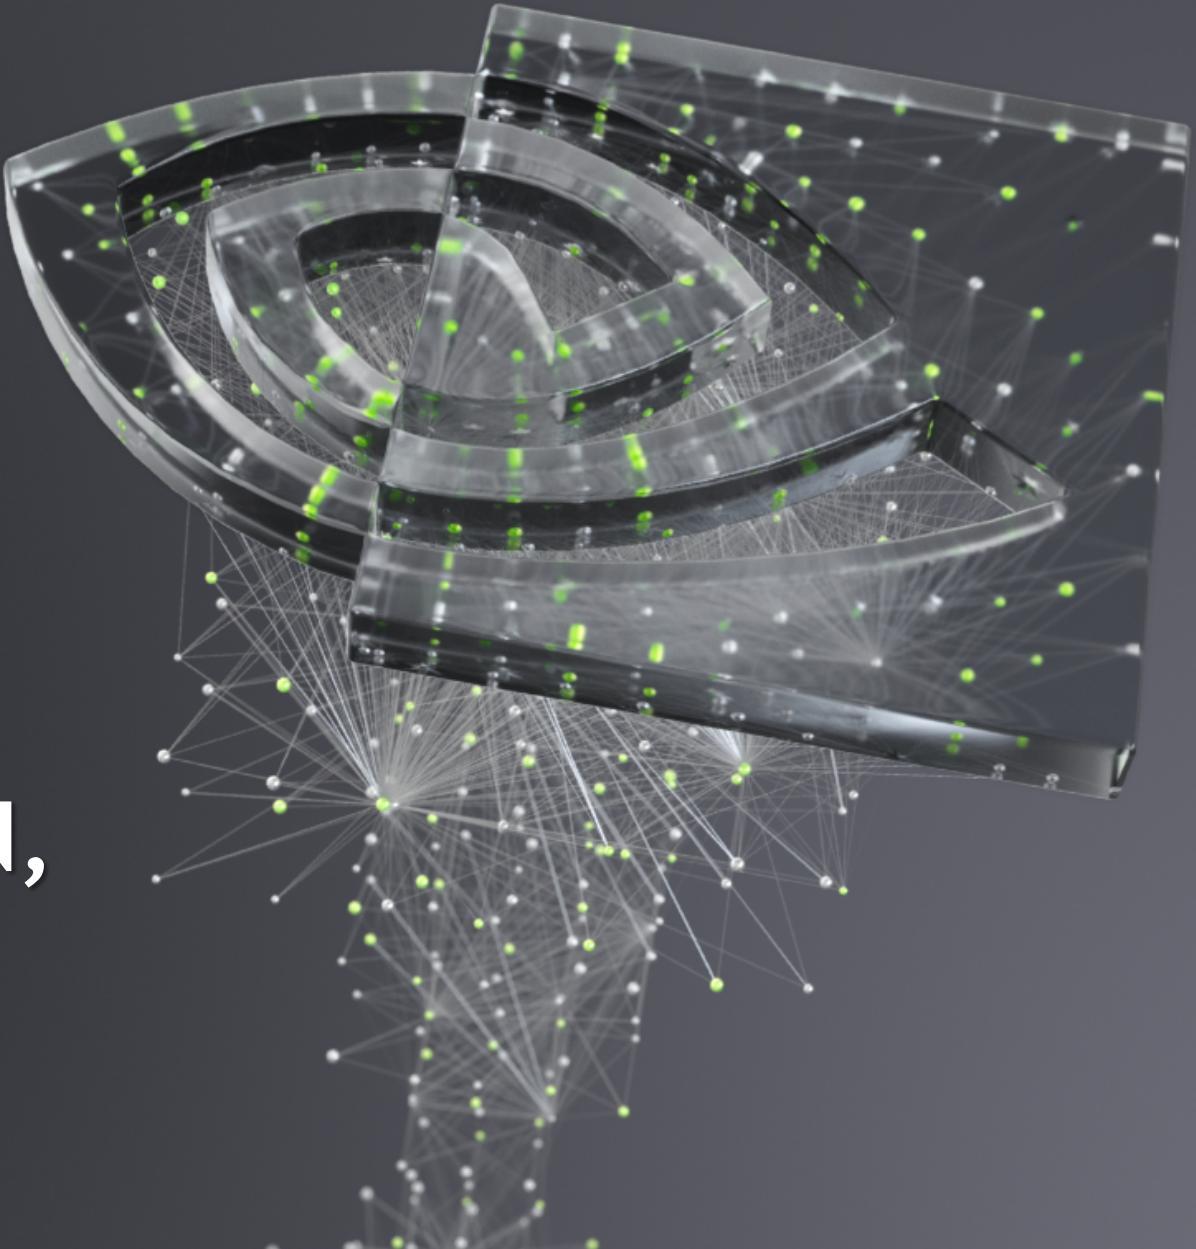




CUDA OPTIMIZATION, PART 2

NVIDIA Corporation



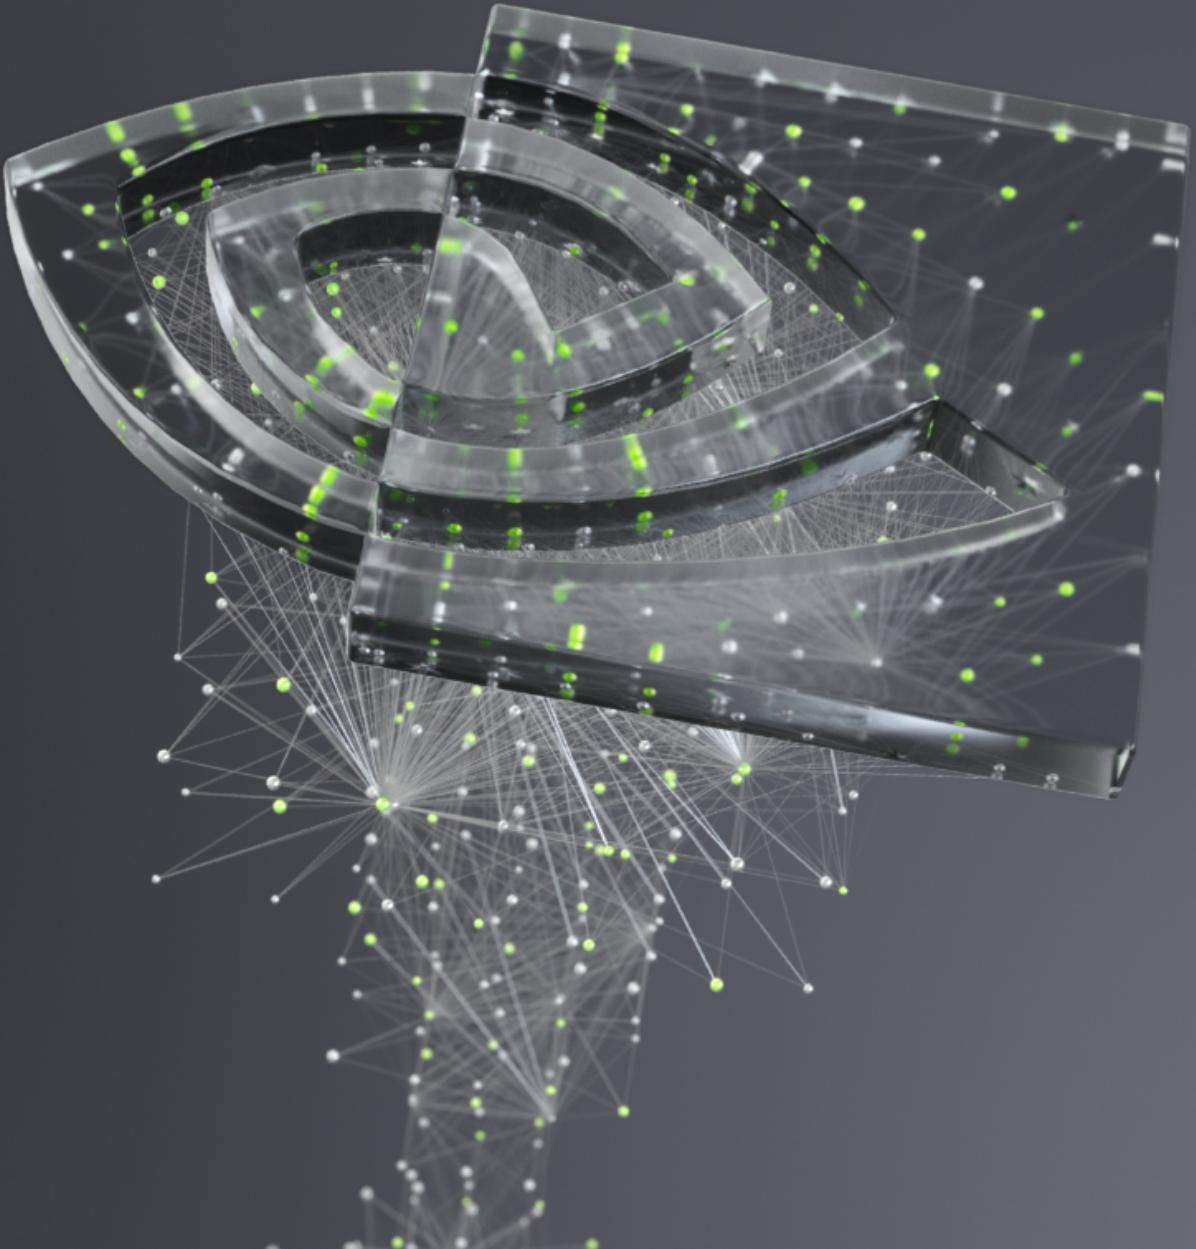
OUTLINE

- ▶ Architecture:
 - Kepler/Maxwell/Pascal/Volta
- ▶ Kernel optimizations
 - ▶ Launch configuration
- ▶ Part 2 (this session):
 - ▶ Global memory throughput
 - ▶ Shared memory access

Most concepts in this presentation apply to *any* language or API on NVIDIA GPUs



GLOBAL MEMORY THROUGHPUT



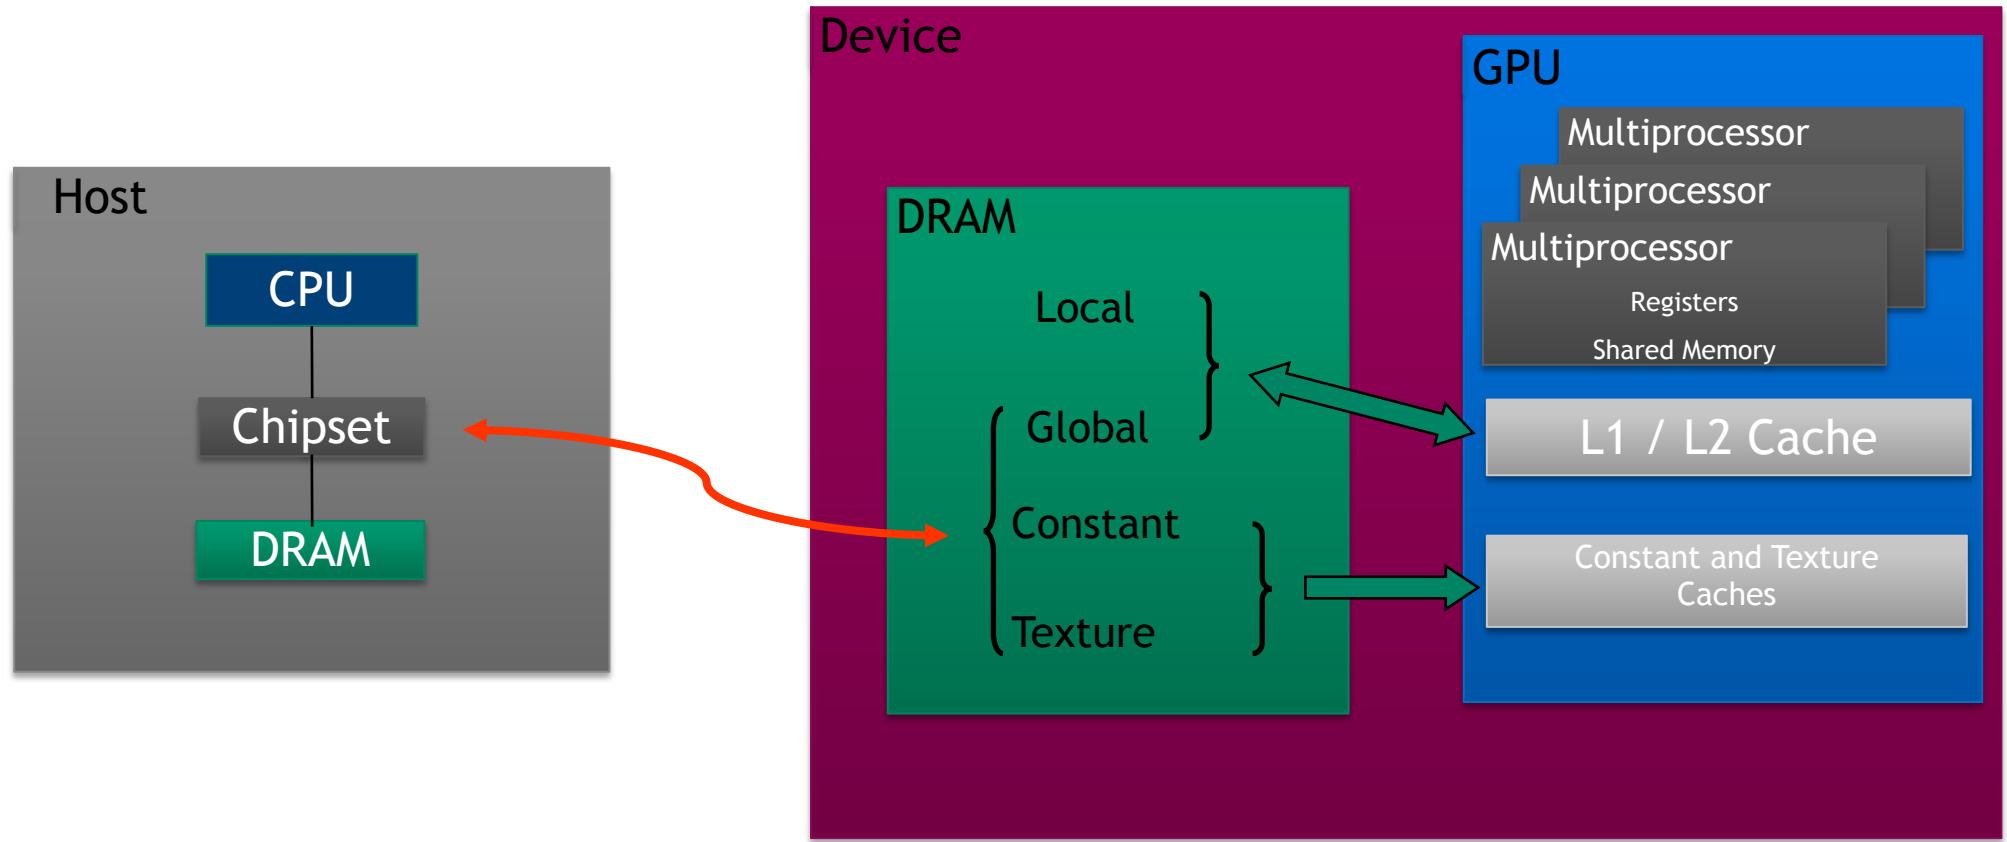
MEMORY HIERARCHY REVIEW

- ▶ Local storage
 - ▶ Each thread has own local storage
 - ▶ Typically registers (managed by the compiler)
- ▶ Shared memory / L1
 - ▶ Program configurable: typically up to 48KB shared (or 64KB, or 96KB...)
 - ▶ Shared memory is accessible by threads in the same threadblock
 - ▶ Very low latency
 - ▶ Very high throughput: >1 TB/s aggregate

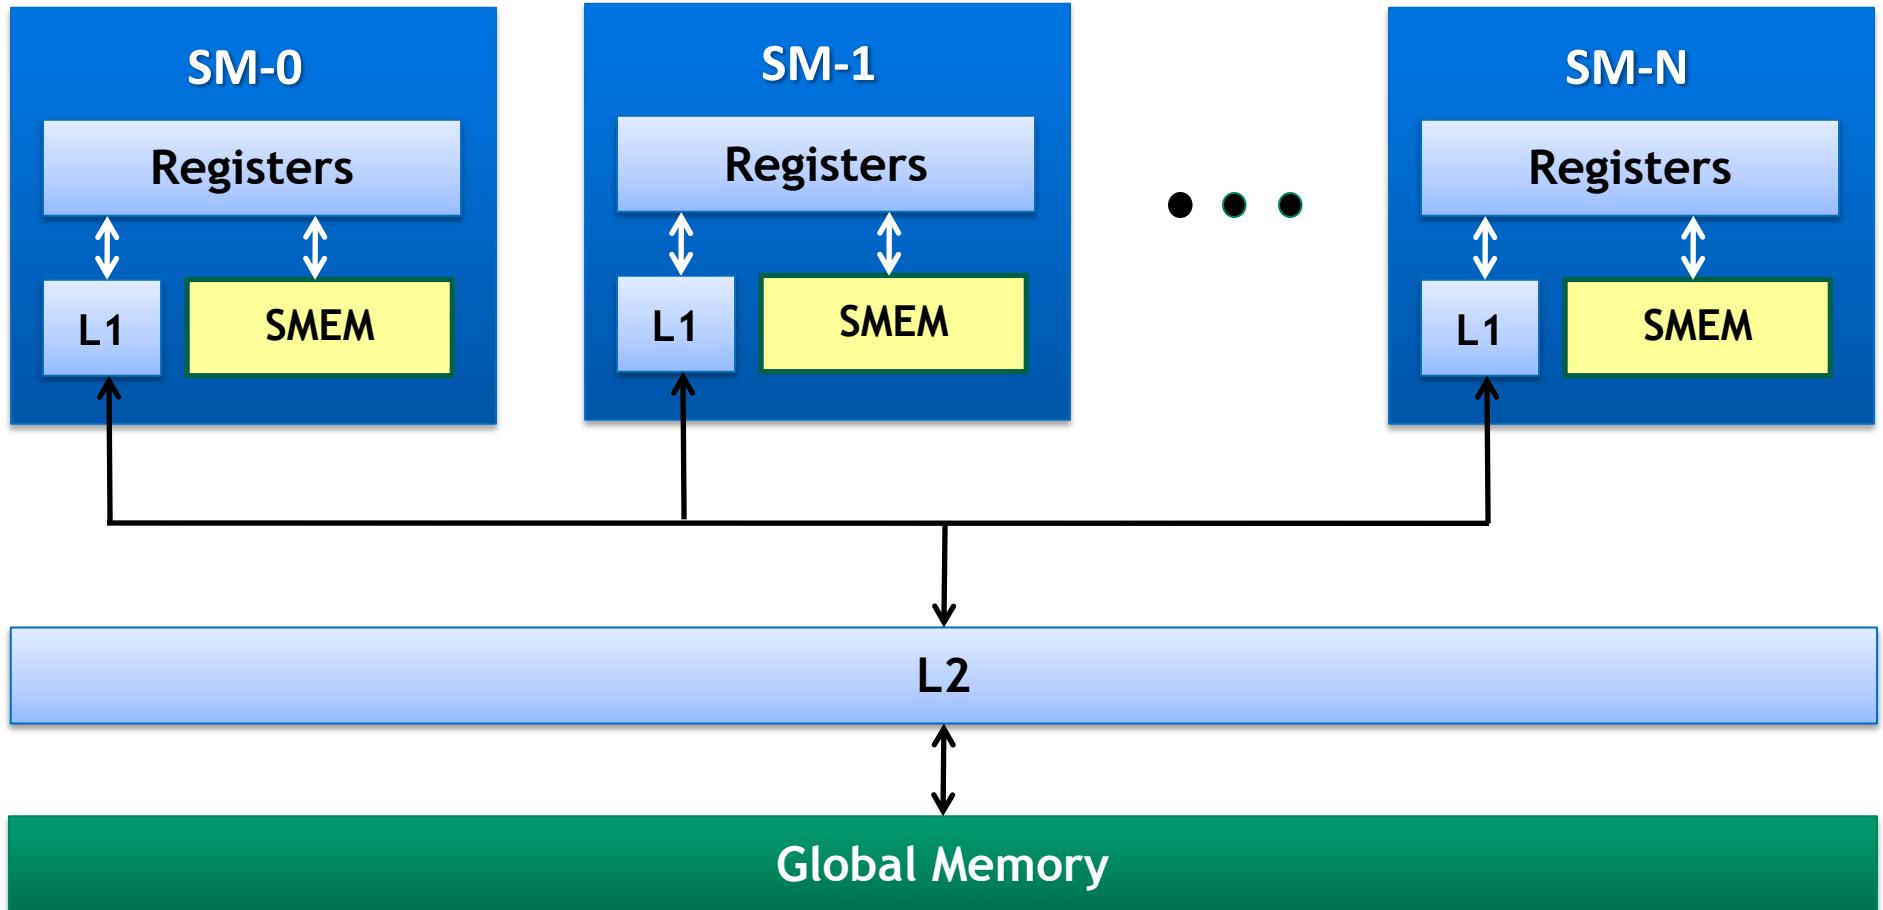
MEMORY HIERARCHY REVIEW

- ▶ L2
 - ▶ All accesses to global memory go through L2, including copies to/from CPU host
- ▶ Global memory
 - ▶ Accessible by all threads as well as host (CPU)
 - ▶ High latency (**hundreds of cycles**)
 - ▶ Throughput: up to **~900 GB/s** (Volta V100)

MEMORY ARCHITECTURE



MEMORY HIERARCHY REVIEW



GMEM OPERATIONS

- ▶ Loads:
 - ▶ Caching
 - ▶ Default mode
 - ▶ Attempts to hit in L1, then L2, then GMEM
 - ▶ Load granularity is **128-byte** line
- ▶ Stores:
 - ▶ Invalidate L1, write-back for L2

GMEM OPERATIONS



- Loads:



- Non-caching

- Compile with `-Xptxas -dlcm=cg` option to nvcc

- Attempts to hit in L2, then GMEM

Do not hit in L1, invalidate the line if it's in L1 already

- Load granularity is 32-bytes

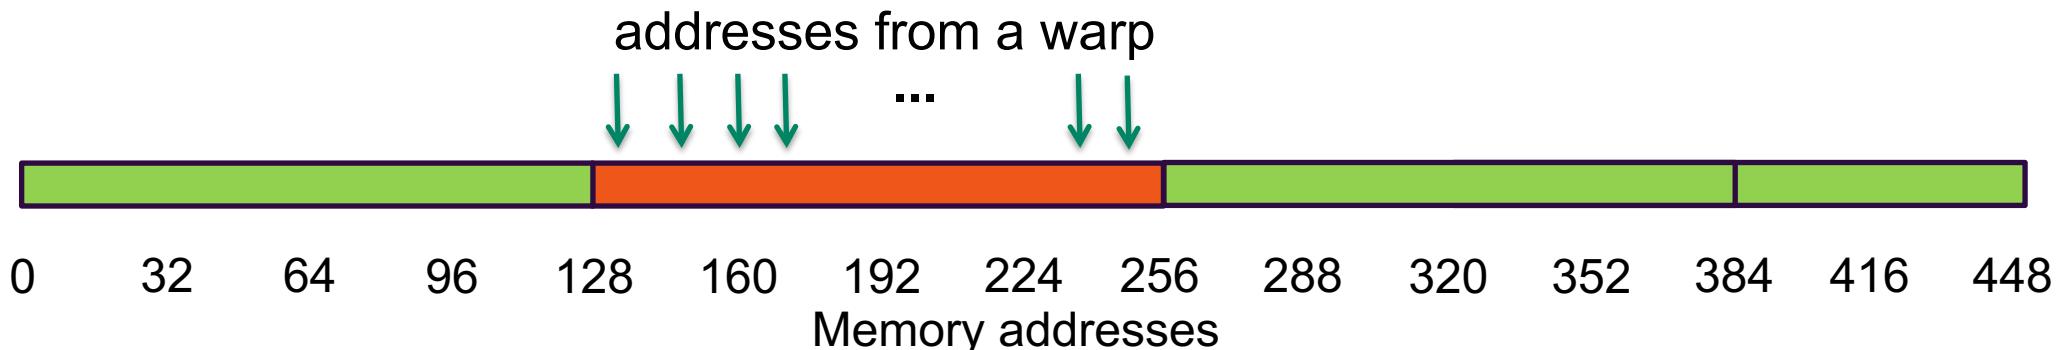
We won't spend much time with non-caching loads in this training session

LOAD OPERATION

- ▶ Memory operations are issued **per warp** (32 threads)
 - ▶ Just like all other instructions
- ▶ Operation:
 - ▶ Threads in a warp provide memory addresses
 - ▶ Determine which lines/segments are needed
 - ▶ Request the needed lines/segments

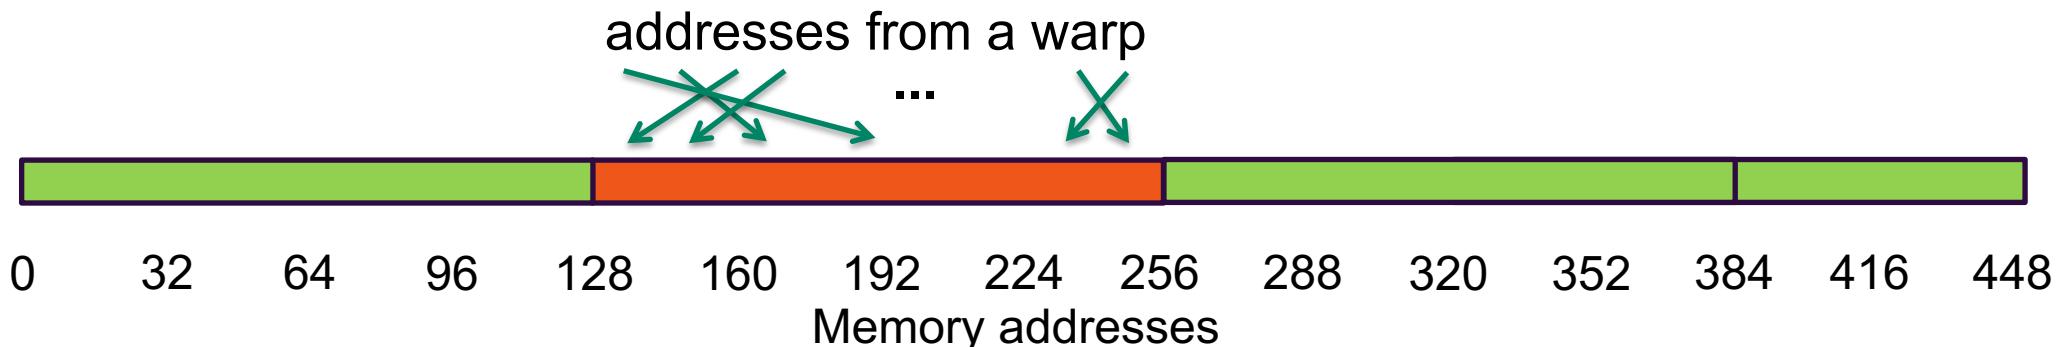
CACHING LOAD

- ▶ Warp requests 32 aligned, consecutive 4-byte words
- ▶ Addresses fall within 1 cache-line
 - ▶ Warp needs 128 bytes
 - ▶ 128 bytes move across the bus on a miss
 - ▶ Bus utilization: 100%
 - ▶ `int c = a[idx];`



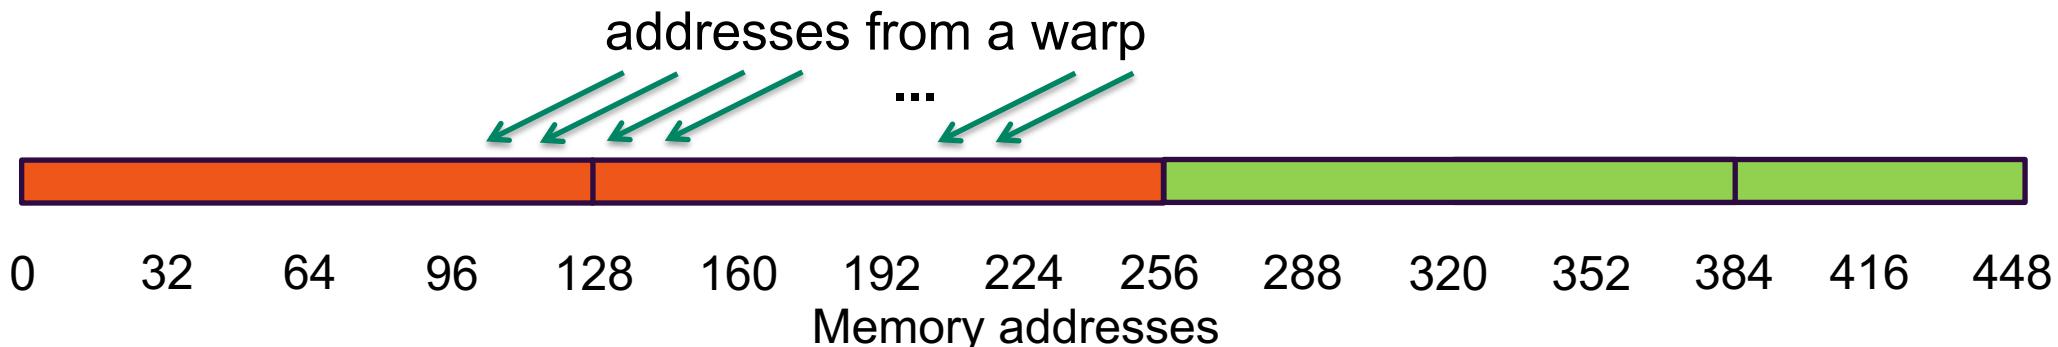
CACHING LOAD

- ▶ Warp requests 32 aligned, permuted 4-byte words
- ▶ Addresses fall within 1 cache-line
 - ▶ Warp needs 128 bytes
 - ▶ 128 bytes move across the bus on a miss
 - ▶ Bus utilization: **100%**
 - ▶ `int c = a[rand()%warpSize];`



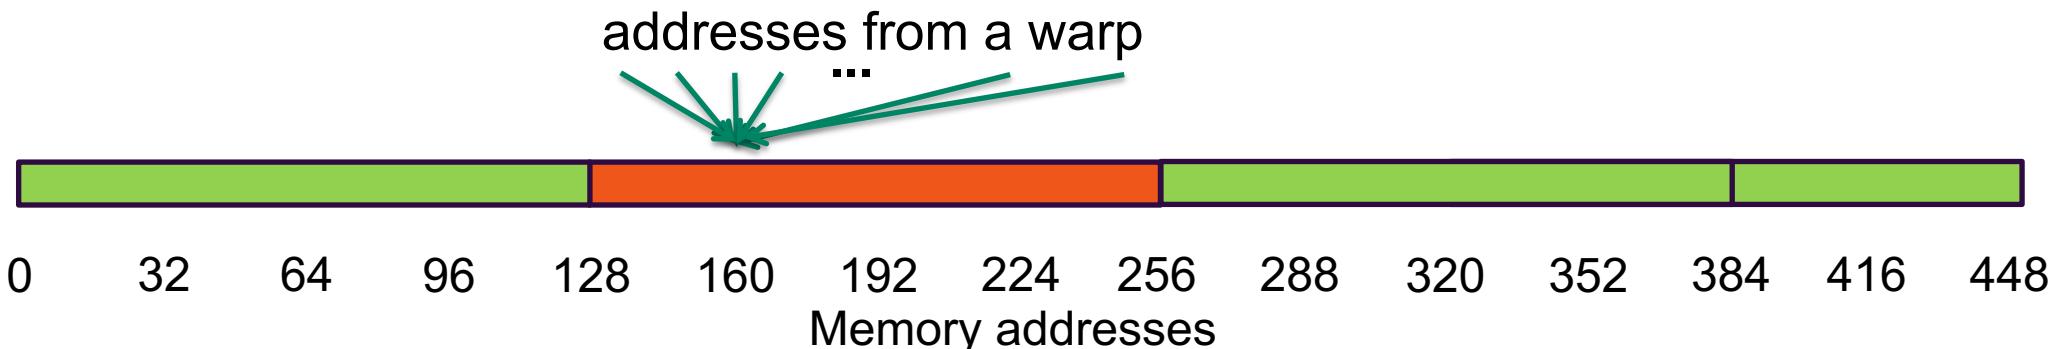
CACHING LOAD

- ▶ Warp requests 32 misaligned, consecutive 4-byte words
- ▶ Addresses fall within 2 cache-lines
 - ▶ Warp needs 128 bytes
 - ▶ 256 bytes move across the bus on misses
 - ▶ Bus utilization: 50%
 - ▶ `int c = a[idx-2];`



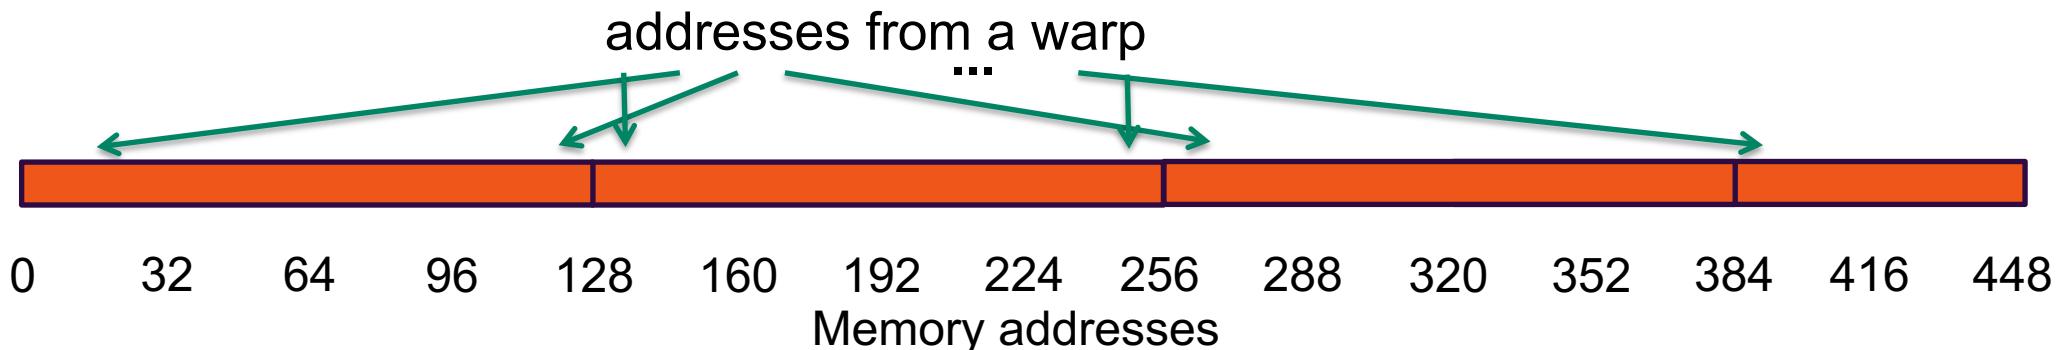
CACHING LOAD

- ▶ All threads in a warp request the same 4-byte word
- ▶ Addresses fall within a single cache-line
 - ▶ Warp needs 4 bytes
 - ▶ 128 bytes move across the bus on a miss
 - ▶ Bus utilization: 3.125%
 - ▶ `int c = a[40];`



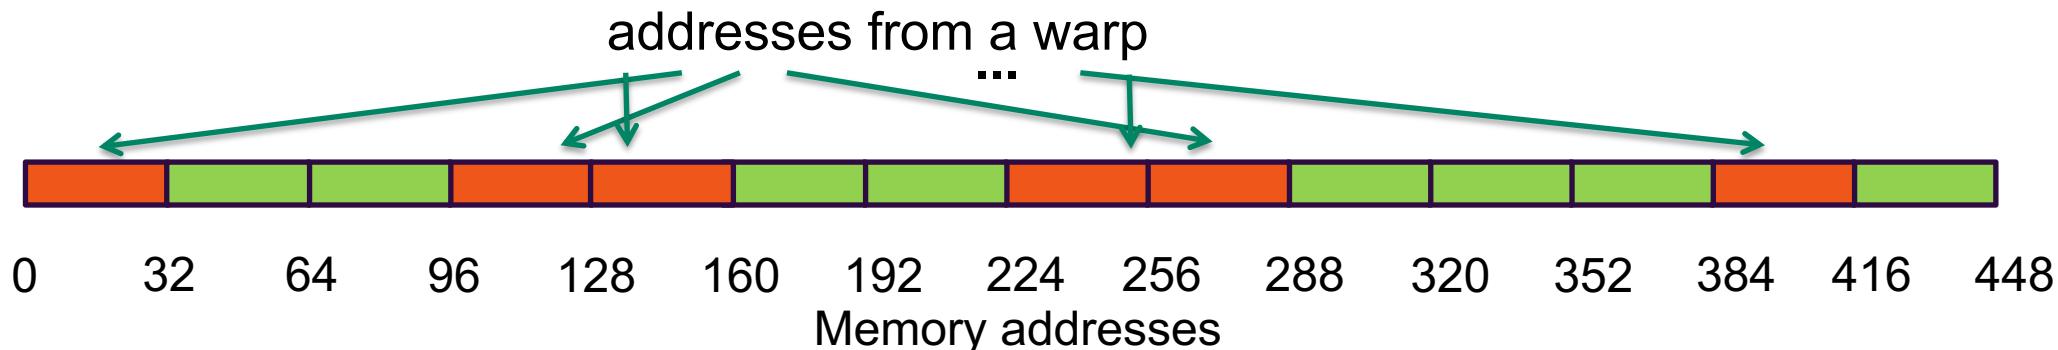
CACHING LOAD

- ▶ Warp requests 32 scattered 4-byte words
- ▶ Addresses fall within N cache-lines
 - ▶ Warp needs 128 bytes
 - ▶ $N \times 128$ bytes move across the bus on a miss
 - ▶ Bus utilization: $128 / (N \times 128)$ (3.125% worst case $N=32$)
 - ▶ `int c = a[rand()];`



NON-CACHING LOAD

- ▶ Warp requests 32 scattered 4-byte words
 - ▶ Addresses fall within N segments
 - ▶ Warp needs 128 bytes
 - ▶ $N \times 32$ bytes move across the bus on a miss
 - ▶ Bus utilization: $128 / (N \times 32)$ (12.5% worst case $N = 32$)
 - ▶ `int c = a[rand()]; -Xptxas -dlcm=cg`

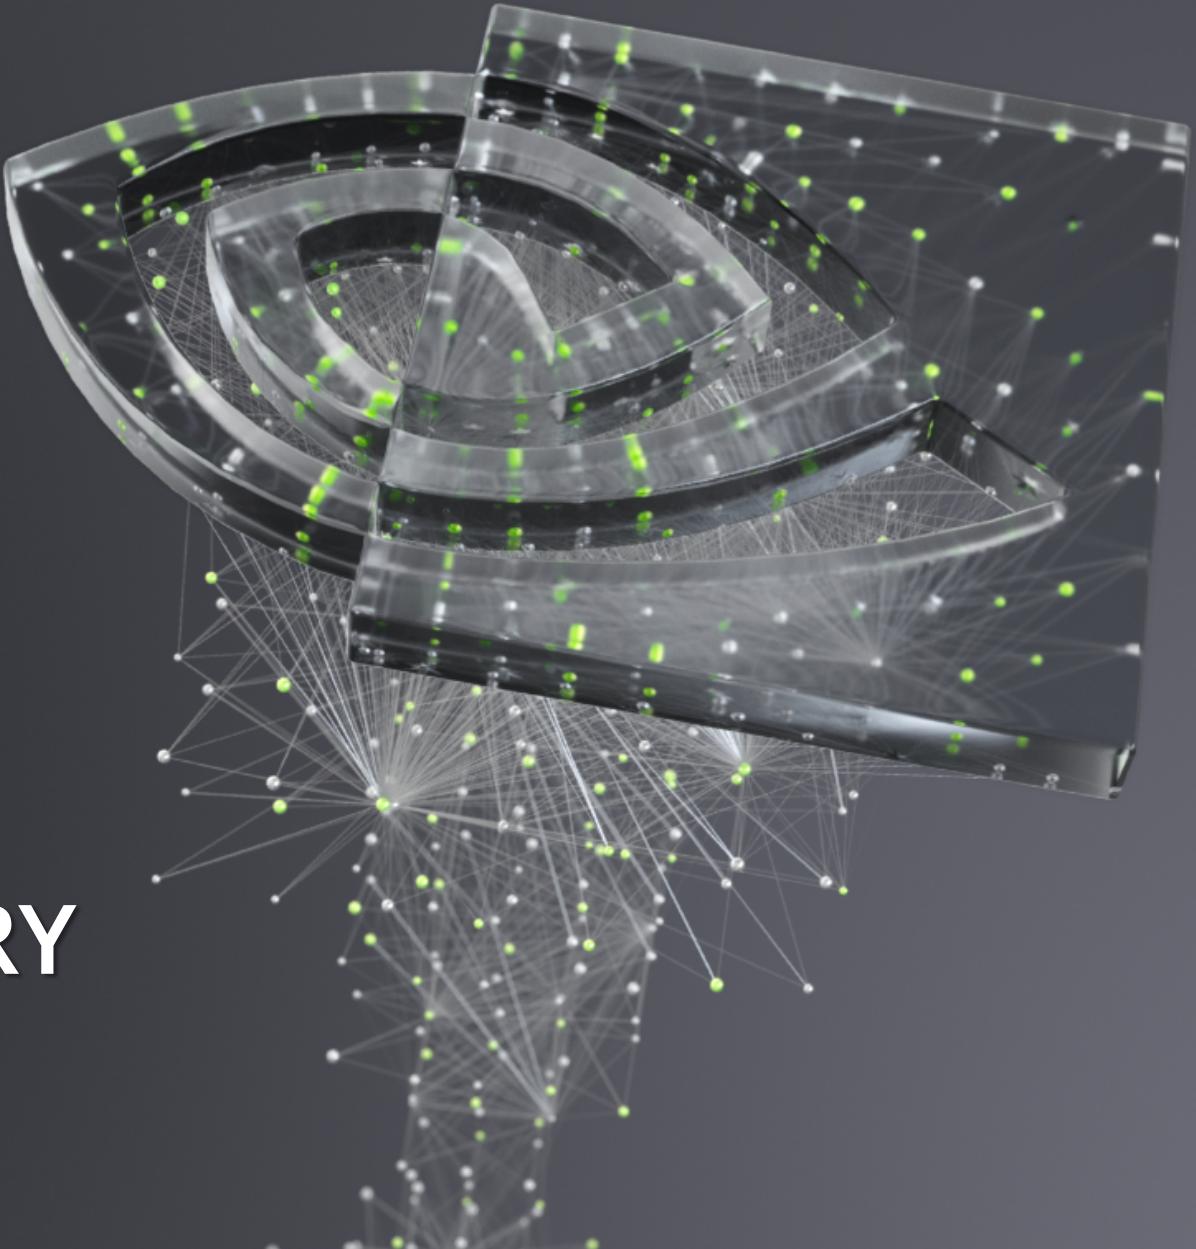


GMEM OPTIMIZATION GUIDELINES

- ▶ Strive for perfect coalescing
 - ▶ (Align starting address - may require padding)
 - ▶ A warp should access within a contiguous region
- ▶ Have enough concurrent accesses to saturate the bus
 - ▶ Process several elements per thread
 - ▶ Multiple loads get pipelined
 - ▶ Indexing calculations can often be reused
 - ▶ Launch enough threads to maximize throughput
 - ▶ Latency is hidden by switching threads (warps)
- ▶ Use all the caches!



SHARED MEMORY



SHARED MEMORY

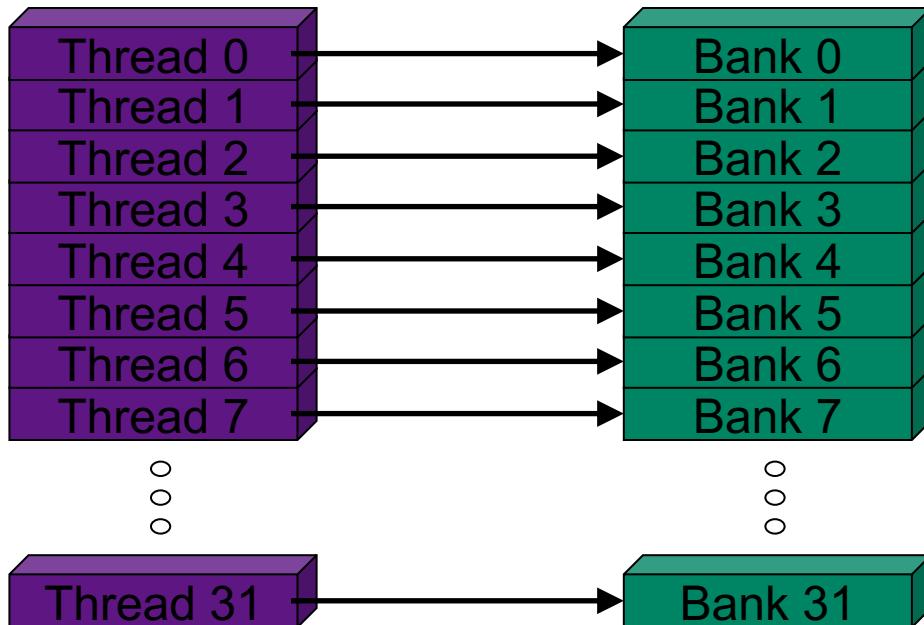
- ▶ Uses:
 - ▶ Inter-thread communication within a block
 - ▶ Cache data to reduce redundant global memory accesses
 - ▶ Use it to improve global memory access patterns
- ▶ Organization:
 - ▶ 32 banks, **4-byte** wide banks
 - ▶ Successive 4-byte words belong to different banks

SHARED MEMORY

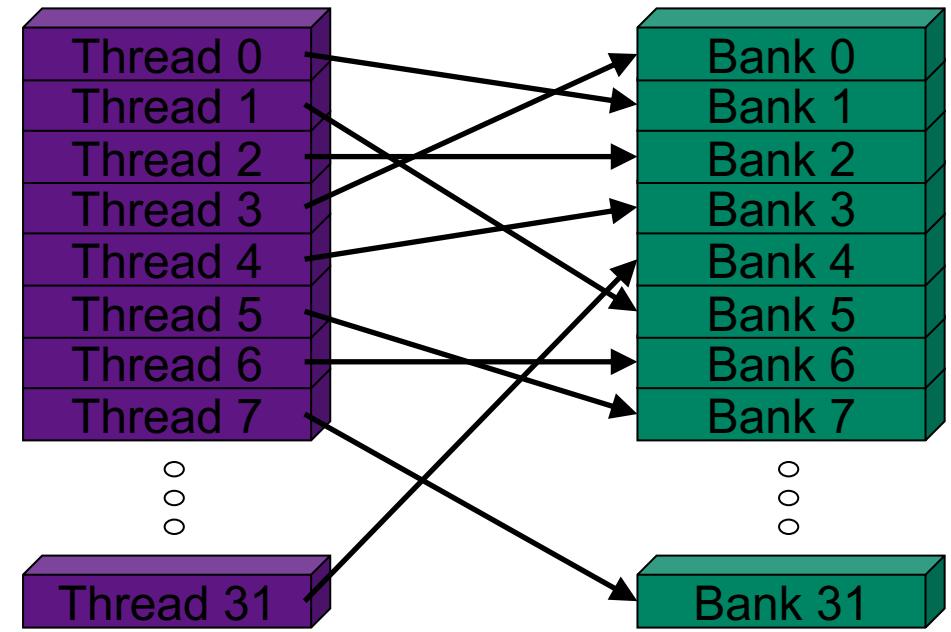
- ▶ Performance:
 - ▶ Typically: 4 bytes per bank per 1 or 2 clocks per multiprocessor
 - ▶ shared accesses are issued per 32 threads (warp)
 - ▶ serialization: if N threads of 32 access different 4-byte words in the same bank, N accesses are executed serially
 - ▶ multicast: N threads access the same word in one fetch
 - ▶ Could be different bytes within the same word

BANK ADDRESSING EXAMPLES

No Bank Conflicts

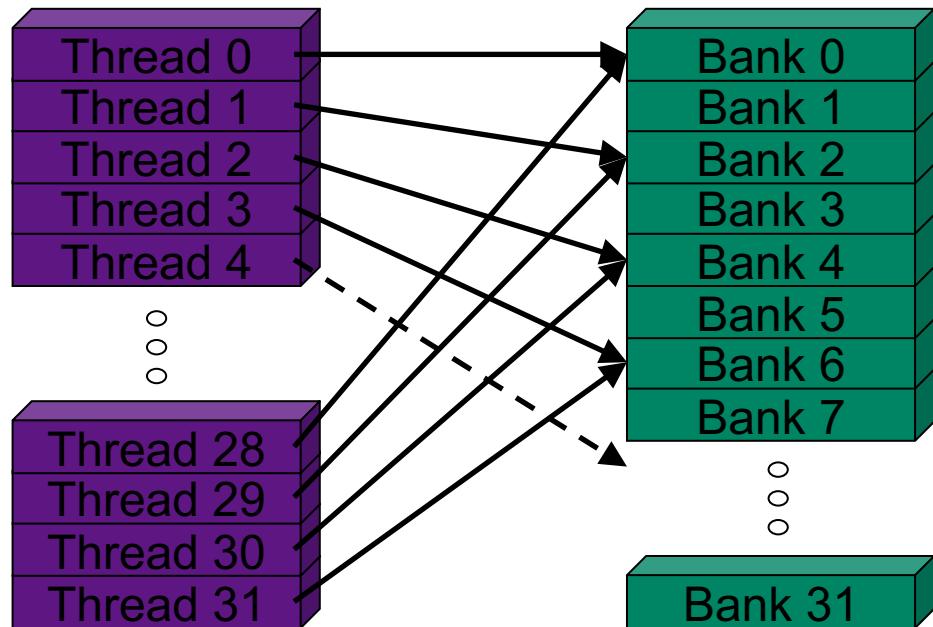


No Bank Conflicts

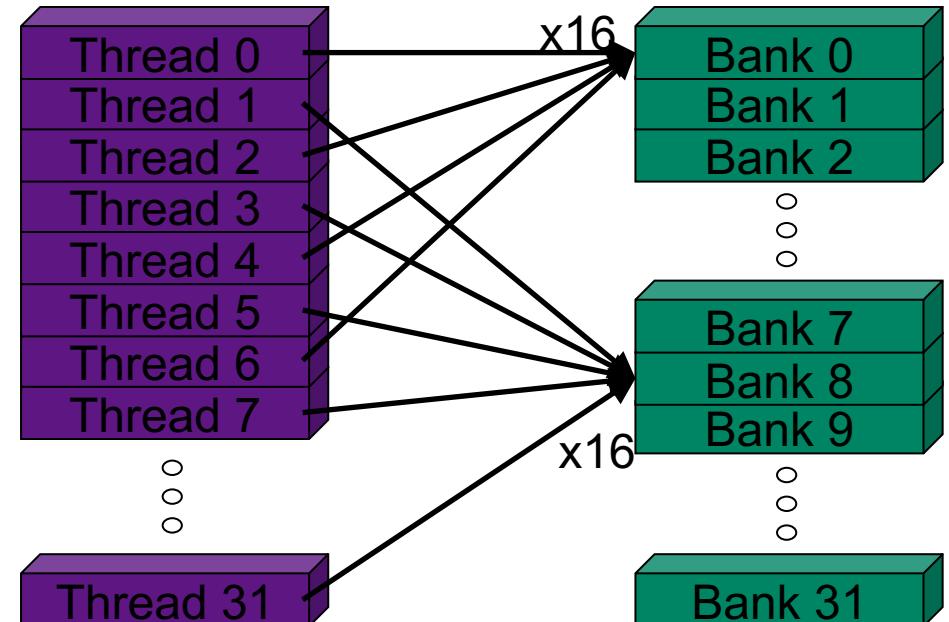


BANK ADDRESSING EXAMPLES

2-way Bank Conflicts

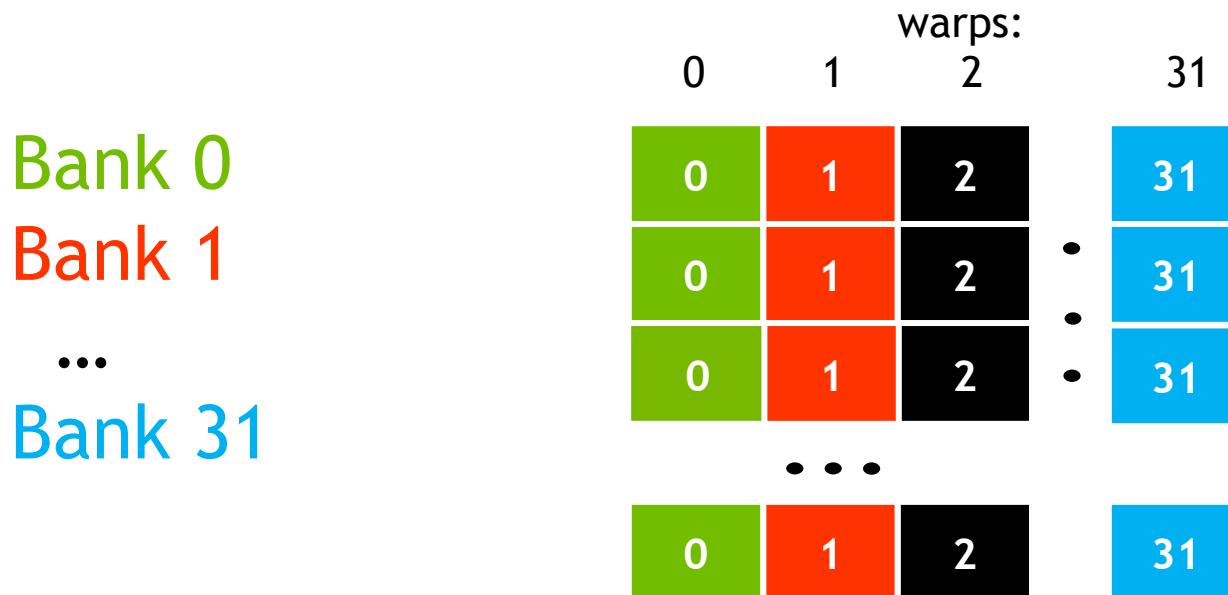


16-way Bank Conflicts



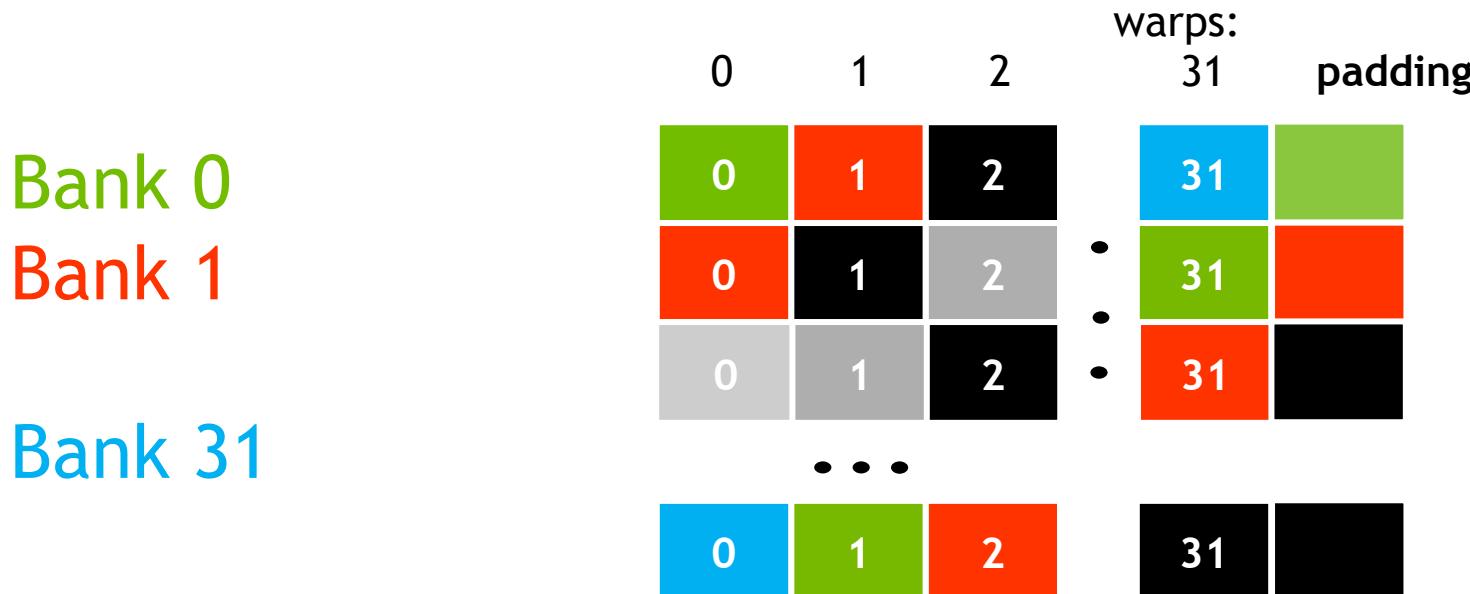
SHARED MEMORY: AVOIDING BANK CONFLICTS

- ▶ 32x32 SMEM array
 - ▶ Warp accesses a column:
 - ▶ 32-way bank conflicts (threads in a warp access the same bank)



SHARED MEMORY: AVOIDING BANK CONFLICTS

- ▶ Add a column for padding:
 - ▶ 32x33 SMEM array
- ▶ Warp accesses a column:
 - ▶ 32 different banks, no bank conflicts



SUMMARY

- ▶ Kernel Launch Configuration:
 - ▶ Launch enough threads per SM to hide latency
 - ▶ Launch enough threadblocks to load the GPU
- ▶ Global memory:
 - ▶ Maximize throughput (GPU has lots of bandwidth, use it effectively)
- ▶ Use shared memory when applicable (over 1 TB/s bandwidth)
- ▶ Use analysis/profiling when optimizing:
 - ▶ “Analysis-driven Optimization” (future session)

FUTURE SESSIONS

- ▶ Atomics, Reductions, Warp Shuffle
- ▶ Using Managed Memory
- ▶ Concurrency (streams, copy/compute overlap, multi-GPU)
- ▶ Analysis Driven Optimization
- ▶ Cooperative Groups

FURTHER STUDY

- ▶ Optimization in-depth:
 - ▶ <http://on-demand.gputechconf.com/gtc/2013/presentations/S3466-Programming-Guidelines-GPU-Architecture.pdf>
- ▶ Analysis-Driven Optimization:
 - ▶ <http://on-demand.gputechconf.com/gtc/2012/presentations/S0514-GTC2012-GPU-Performance-Analysis.pdf>
- ▶ CUDA Best Practices Guide:
 - ▶ <https://docs.nvidia.com/cuda/cuda-c-best-practices-guide/index.html>
- ▶ CUDA Tuning Guides:
 - ▶ <https://docs.nvidia.com/cuda/index.html#programming-guides>
(Kepler/Maxwell/Pascal/Volta)

HOMEWORK

- ▶ Log into Summit (ssh username@home.ccs.ornl.gov -> ssh summit)
- ▶ Clone GitHub repository:
 - ▶ Git clone [git@github.com:olcf/cuda-training-series.git](https://github.com/olcf/cuda-training-series.git)
- ▶ Follow the instructions in the readme.md file:
 - ▶ <https://github.com/olcf/cuda-training-series/blob/master/exercises/hw4/readme.md>
- ▶ Prerequisites: basic linux skills, e.g. ls, cd, etc., knowledge of a text editor like vi/emacs, and some knowledge of C/C++ programming



QUESTIONS?

